

# The Case for Writing Network Drivers in High-Level Programming Languages

**Paul Emmerich, Simon Ellmann**, Fabian Bonk, Alex Egger,  
Esaú García Sánchez-Torija, Thomas Günzel, Sebastian Di Luzio,  
Alexandru Obada, Maximilian Stadlmeier, Sebastian Voit, Georg Carle

September 24, 2019

Chair of Network Architectures and Services  
Department of Informatics  
Technical University of Munich

# C is an awesome language for operating systems!

- Low-level access to memory and devices
- Pointers are awesome
- Everyone can read and write C
- You can write safe and secure code if you try really hard

# C can cause security problems

Vulnerability Trends Over Time

Year	# of Vulnerabilities	DoS	Code Execution	Overflow	Memory Corruption	Sql Injection	XSS	Directory Traversal	Http Response Splitting	Bypass something	Gain Information	Gain Privileges
<a href="#">1999</a>	19	<a href="#">2</a>		<a href="#">3</a>						<a href="#">1</a>		<a href="#">2</a>
<a href="#">2000</a>	5	<a href="#">3</a>										<a href="#">1</a>
<a href="#">2001</a>	22	<a href="#">6</a>								<a href="#">4</a>		<a href="#">3</a>
<a href="#">2002</a>	15	<a href="#">3</a>		<a href="#">1</a>						<a href="#">1</a>	<a href="#">1</a>	
<a href="#">2003</a>	19	<a href="#">8</a>		<a href="#">2</a>						<a href="#">1</a>	<a href="#">3</a>	<a href="#">4</a>
<a href="#">2004</a>	51	<a href="#">20</a>	<a href="#">5</a>	<a href="#">12</a>							<a href="#">5</a>	<a href="#">12</a>

(...)

<a href="#">2017</a>	454	<a href="#">147</a>	<a href="#">169</a>	<a href="#">52</a>	<a href="#">26</a>			<a href="#">1</a>		<a href="#">17</a>	<a href="#">89</a>	<a href="#">36</a>
<a href="#">2018</a>	166	<a href="#">81</a>	<a href="#">3</a>	<a href="#">28</a>	<a href="#">8</a>					<a href="#">3</a>	<a href="#">17</a>	<a href="#">3</a>
Total	2155	<a href="#">1184</a>	<a href="#">241</a>	<a href="#">347</a>	<a href="#">124</a>			<a href="#">3</a>		<a href="#">111</a>	<a href="#">350</a>	<a href="#">260</a>
% Of All		54.9	11.2	16.1	5.8	0.0	0.0	0.1	0.0	5.2	16.2	12.1

- Screenshot from <https://www.cvedetails.com/>
- Security bugs found in the Linux kernel in the last  $\approx 20$  years

# C can cause security problems

- Not all bugs can be blamed on the language
- Cutler et al. analyzed 65 CVEs categorized as code execution in the Linux kernel <sup>1</sup>

---

<sup>1</sup> C. Cutler, M. F. Kaashoek, and R. T. Morris, “[The benefits and costs of writing a POSIX kernel in a high-level language](#)”, USENIX OSDI, 2018

## C can cause security problems

- Not all bugs can be blamed on the language, but 61% can
- Cutler et al. analyzed 65 CVEs categorized as code execution in the Linux kernel <sup>1</sup>

Bug type	Num.	Perc.	Can be avoided by using a better language?
Various	11	17%	Unclear/Maybe
Logic	14	22%	No
Use-after-free	8	12%	Yes
Out of bounds	32	49%	Yes (likely leads to panic)

**Table 1:** Code execution vulnerabilities in the Linux kernel identified by Cutler et al.<sup>1</sup>

<sup>1</sup> C. Cutler, M. F. Kaashoek, and R. T. Morris, “[The benefits and costs of writing a POSIX kernel in a high-level language](#)”, USENIX OSDI, 2018

# Let's rewrite all operating systems in better languages?

- Rewriting the whole operating system in a safer language is a laudable effort
  - Redox (Rust) wants to become a production-grade OS but currently isn't
  - Singularity (Sing#, Microsoft Research) demonstrated some interesting concepts
  - Biscuit (Go) implements parts of POSIX for research
  - Unikernels like MirageOS (OCaml) or IncludeOS (C++) can be useful in some scenarios

# Let's rewrite all operating systems in better languages?

- Rewriting the whole operating system in a safer language is a laudable effort
  - Redox (Rust) wants to become a production-grade OS but currently isn't
  - Singularity (Sing#, Microsoft Research) demonstrated some interesting concepts
  - Biscuit (Go) implements parts of POSIX for research
  - Unikernels like MirageOS (OCaml) or IncludeOS (C++) can be useful in some scenarios
- But none of these will replace your main operating system any time soon

# Where are these bugs that could have been prevented?

- We looked at these 40 preventable bugs
- 39 of them were in drivers (the other was in the Bluetooth stack)



# Where are these bugs that could have been prevented?

- We looked at these 40 preventable bugs
- 39 of them were in drivers (the other was in the Bluetooth stack)
- 13 were in the Qualcomm WiFi driver

# Can we rewrite drivers in better languages?

- User space drivers can be written in **any** language!
- But are all languages an equally good choice?
- Is a JIT compiler or a garbage collector a problem in a driver?

# Challenges for high-level languages

- Access to `mmap` with the proper flags
- Handle externally allocated (foreign) memory in the language
- Handle memory layouts/formats (i.e., access memory that looks like a given C struct)
- Memory access semantics: memory barriers, volatile reads/writes
- Some operations in drivers are inherently unsafe

# Why look at network drivers?

- Easy to benchmark to quantify results
- Huge attack surface: exposed to the external world by design
- User space network drivers are already quite common (e.g., DPDK, Snabb)
- Network stacks are also moving into the user space (e.g., QUIC)

# Why look at network drivers?

- Easy to benchmark to quantify results
- Huge attack surface: exposed to the external world by design
- User space network drivers are already quite common (e.g., DPDK, Snabb)
- Network stacks are also moving into the user space (e.g., QUIC)
- Everything mentioned here is applicable to other drivers as well

A scatter plot showing the relationship between the maximum supported speed and the number of lines of code for DPDK and Linux drivers. The x-axis represents 'Max supported speed' with categories: 10M, 100M, 1G, 2.5G, 10G, 40G, and 100G. The y-axis represents 'Lines of code' on a logarithmic scale from  $10^2$  to  $10^5$ . DPDK drivers are marked with blue '+' and Linux drivers with orange 'x'. A black curve represents a polynomial fit to the data, with the equation  $0.3624x + 5781$  displayed on the plot.

Max supported speed	DPDK drivers (Lines of code)	Linux drivers (Lines of code)
10M	-	~600, ~1,500, ~3,000
100M	-	~400, ~800, ~1,000, ~1,500, ~2,000, ~3,000, ~4,000, ~6,000, ~10,000
1G	-	~150, ~600, ~800, ~1,000, ~1,200, ~1,500, ~2,000, ~3,000, ~4,000, ~6,000, ~10,000, ~20,000
2.5G	~1,000	~5,000
10G	~2,500, ~3,500, ~5,000, ~6,000, ~7,000, ~8,000, ~10,000, ~15,000, ~25,000, ~40,000	~1,500, ~2,000, ~3,000, ~4,000, ~5,000, ~6,000, ~8,000, ~10,000, ~15,000, ~20,000, ~30,000, ~40,000, ~50,000, ~60,000, ~80,000
40G	~8,000, ~10,000, ~15,000, ~20,000, ~25,000, ~30,000, ~40,000, ~50,000, ~60,000, ~70,000, ~80,000, ~90,000	~8,000, ~10,000, ~15,000, ~20,000, ~25,000, ~30,000, ~40,000, ~50,000, ~60,000, ~70,000, ~80,000, ~90,000, ~100,000
100G	~15,000, ~20,000, ~25,000, ~30,000, ~40,000, ~50,000, ~60,000, ~70,000, ~80,000, ~90,000, ~100,000	~8,000, ~10,000, ~15,000, ~20,000, ~25,000, ~30,000, ~40,000, ~50,000, ~60,000, ~70,000, ~80,000, ~90,000, ~100,000

We wrote full user space network drivers in these languages

C#



Swift



OCaml



# Goals for our implementations

- Implement the same feature set as our ixy C driver
- Use a similar structure and architecture as ixy
- Write idiomatic code for the selected language
- Use language safety features where possible
- Quantify trade-offs for performance vs. safety



## Language comparison: Safety properties

Language	General memory		Packet buffers		Int overflows
	Bounds checks	Use after free	Bounds checks	Use after free	
C	X	X	X	X	X
Rust					
Go					
C#					
Java					
OCaml					
Haskell					
Swift					
JavaScript					
Python					

Table 2: Language-level protections against classes of bugs in our drivers

# Language comparison: Safety properties

Language	General memory		Packet buffers		Int overflows
	Bounds checks	Use after free	Bounds checks	Use after free	
C	✗	✗	✗	✗	✗
Rust	✓	✓	(✓) <sup>1</sup>	✓	(✓) <sup>4</sup>
Go	✓	✓	(✓) <sup>1</sup>	(✓) <sup>3</sup>	✗
C#	✓	✓	(✓) <sup>1</sup>	(✓) <sup>3</sup>	(✓) <sup>4</sup>
Java	✓	✓	(✓) <sup>1</sup>	(✓) <sup>3</sup>	✗
OCaml	✓	✓	(✓) <sup>1</sup>	(✓) <sup>3</sup>	✗
Haskell	✓	✓	(✓) <sup>1</sup>	(✓) <sup>3</sup>	(✓) <sup>5</sup>
Swift	✓	✓	✗ <sup>2</sup>	(✓) <sup>3</sup>	✓
JavaScript	✓	✓	(✓) <sup>1</sup>	(✓) <sup>3</sup>	(✓) <sup>5</sup>
Python	✓	✓	(✓) <sup>1</sup>	(✓) <sup>3</sup>	(✓) <sup>5</sup>

<sup>1</sup> Bounds enforced by wrapper, constructor in unsafe or C code


<sup>2</sup> Bounds only enforced in debug mode

<sup>3</sup> Buffers are never free'd/gc'd, only returned to a memory pool


<sup>4</sup> Disabled by default

<sup>5</sup> Uses floating point or arbitrary precision integers by default

## Performance comparison: Test setup



## Batching at 3.3 GHz CPU speed (single core)



## Why is Rust slower than C?

Events per packet	Batch 32, 1.6 GHz		Batch 8, 1.6 GHz	
	C	Rust	C	Rust
<b>Cycles</b>	94	100	108	120
<b>Instructions</b>	127	209	139	232
<b>Instr. per cycle</b>	1.35	2.09	1.29	1.93
<b>Branches</b>	18	24	19	27
<b>Branch mispredicts</b>	0.05	0.08	0.02	0.06
<b>Store <math>\mu</math>ops</b>	21.8	37.4	24.4	43.0
<b>Load <math>\mu</math>ops</b>	30.1	77.0	33.4	84.2
<b>Load L1 hits</b>	24.3	75.9	28.8	83.1
<b>Load L2 hits</b>	1.1	0.05	1.2	0.1
<b>Load L3 hits</b>	0.9	0.0	0.5	0.0
<b>Load L3 misses</b>	0.3	0.1	0.3	0.3


**Table 4:** Performance counter readings in events per packet when forwarding packets

## Why is Rust slower than C?


Events per packet	Batch 32, 1.6 GHz		Batch 8, 1.6 GHz	
	C	Rust	C	Rust
<b>Cycles</b>	94	100	108	120
<b>Instructions</b>	127	209	139	232
<b>Instr. per cycle</b>	1.35	2.09	1.29	1.93
<b>Branches</b>	18	24	19	27
<b>Branch mispredicts</b>	0.05	0.08	0.02	0.06
<b>Store <math>\mu</math>ops</b>	21.8	37.4	24.4	43.0
<b>Load <math>\mu</math>ops</b>	30.1	77.0	33.4	84.2
<b>Load L1 hits</b>	24.3	75.9	28.8	83.1
<b>Load L2 hits</b>	1.1	0.05	1.2	0.1
<b>Load L3 hits</b>	0.9	0.0	0.5	0.0
<b>Load L3 misses</b>	0.3	0.1	0.3	0.3

**Table 5:** Performance counter readings in events per packet when forwarding packets

## Tail latency at 1 Mpps




## Tail latency at 1 Mpps







## Tail latency at 1 Mpps




## Tail latency at 1 Mpps



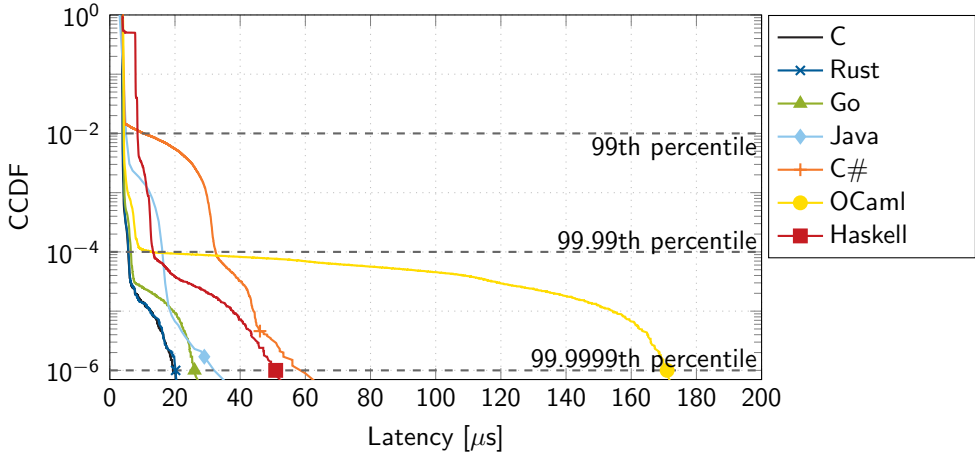
## Tail latency at 1 Mpps



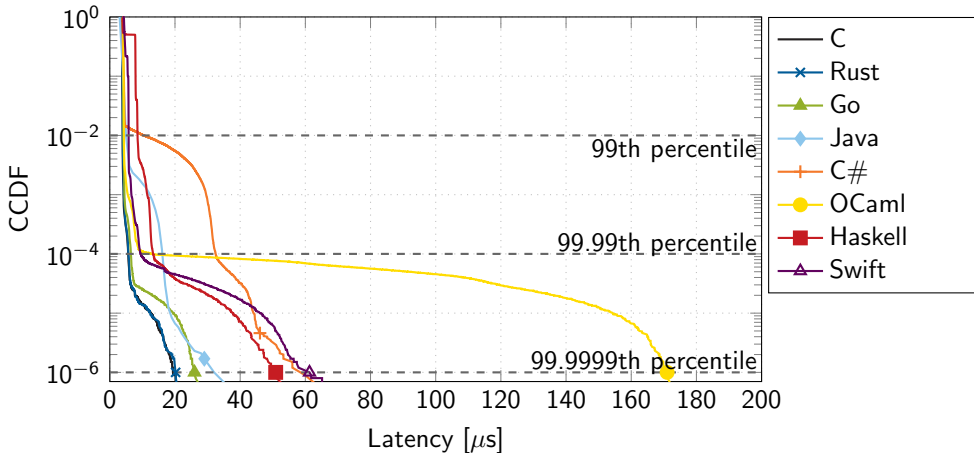
## Tail latency at 1 Mpps




## Tail latency at 1 Mpps



## Tail latency at 1 Mpps



## Tail latency at 1 Mpps



## Conclusion: Check out our code

- Meta-repository with links: <https://github.com/ixy-languages/ixy-languages>
- Should your driver really be in the kernel?
- Next time you write a driver: consider a user space driver in a cool language
- Other cool stuff in the paper: details on implementations, latency at higher loads, Java garbage collector comparison, analysis of user space packet processing frameworks used in academia, study of mistakes made in C, and more...



# Backup Slides

# Languages for code in trustworthy systems

- Rust
  - Fast, no garbage collector
  - Low-level: Easy to reason about performance
  - Safest language of the evaluated languages
- Go
  - Fast, low-latency garbage collector
  - Garbage collector tuned for sub-millisecond latency
  - Easier and faster to write than Rust

# Languages for code in trustworthy systems

- Rust
  - Fast, no garbage collector
  - Low-level: Easy to reason about performance
  - Safest language of the evaluated languages
- Go
  - Fast, low-latency garbage collector
  - Garbage collector tuned for sub-millisecond latency
  - Easier and faster to write than Rust
- Other languages
  - Implement critical parts in different languages in redundant systems
  - Functional languages for easier formal verification

# Language comparison: Overview

Language	Main paradigm	Memory management	Compilation
C	Imperative	No	Compiled
Rust	Imperative	Ownership/RAII	(LLVM) Compiled
Go	Imperative	Garbage collection	Compiled
C#	Object-oriented	Garbage collection	JIT
Java	Object-oriented	Garbage collection	JIT
OCaml	Functional	Garbage collection	Compiled
Haskell	Functional	Garbage collection	(LLVM) Compiled
Swift	Protocol-oriented	Reference counting	(LLVM) Compiled
JavaScript	Imperative	Garbage collection	JIT
Python	Imperative	Garbage collection	Interpreted

Table 6: Language overview

## Language comparison: Implementation sizes

Lang.	Lines of code <sup>1</sup>	Lines of C code <sup>1</sup>	Source size (gzip <sup>2</sup> )
C	831	831	12.9 kB
Rust	961	0	10.4 kB
Go	1640	0	20.6 kB
C#	1266	34	13.1 kB
Java	2885	188	31.8 kB
OCaml	1177	28	12.3 kB
Haskell	1001	0	9.6 kB
Swift	1506	0	15.9 kB
JavaScript	1004	262	13.0 kB
Python	1242	(Cython) 77	14.2 kB

<sup>1</sup> Incl. C code, excluding empty lines and comments, counted with `cloc`

<sup>2</sup> Compression level 6

**Table 7:** Size of our implementations (w/o register constants, stripped features not found in all drivers)